Big Data Hadoop 'Session 9:Advance Hive Assignment 1'

DATASET

This Data set is about Olympics. You can download the data set from the below link:

https://drive.google.com/open?id=0ByJLBTmJojjzV1czX3Nha0R3bTQ

DATE SET DESCRIPTION

The data set consists of the following fields

ete: This field consists of the athlete name
This field consists of athlete ages
ntry: This fields consists of the country names which participated in Olympics
: This field consists of the year
ing Date: This field consists of the closing date of ceremony
t: Consists of the sports name
Medals: No. of Gold medals
r Medals: No. of Silver medals
ze Medals: No. of Bronze medals
l Medals: Consists of total no. of medals

LOAD DATA in Hive olympic table

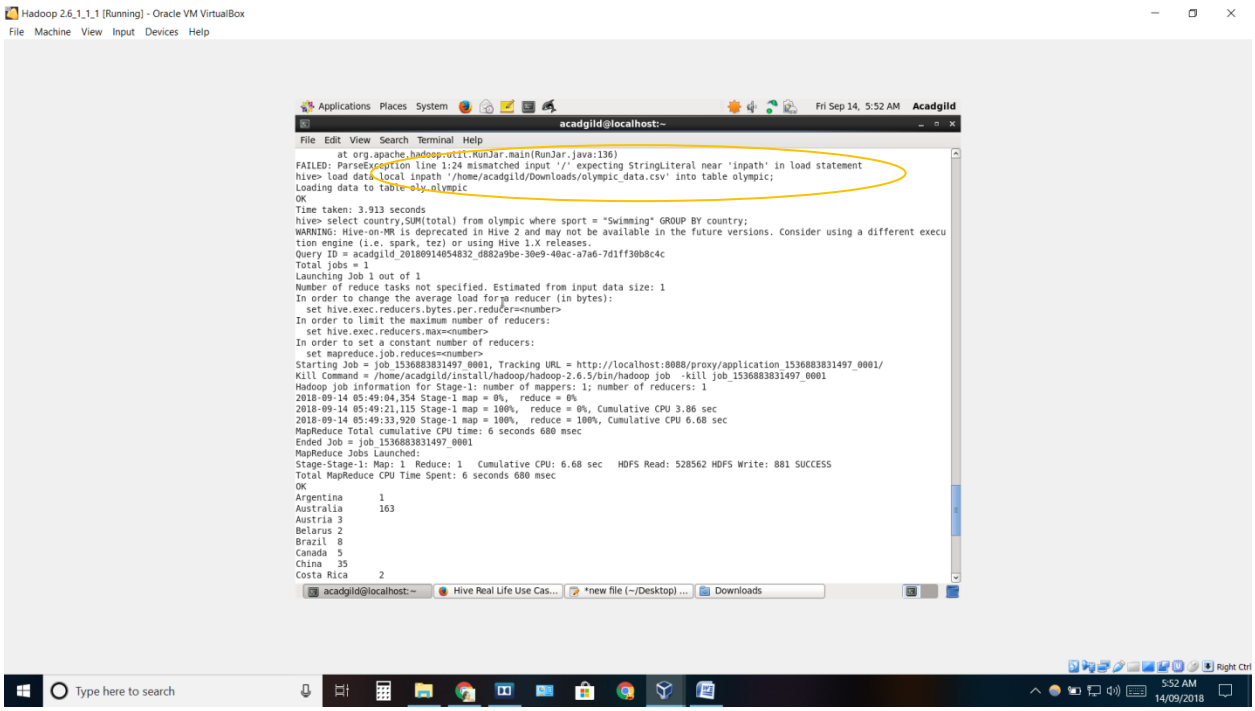

Task 1

1. Write a Hive program to find the number of medals won by each country in swimming.

   select country,SUM(total) from olympic where sport = "Swimming" GROUP BY country;

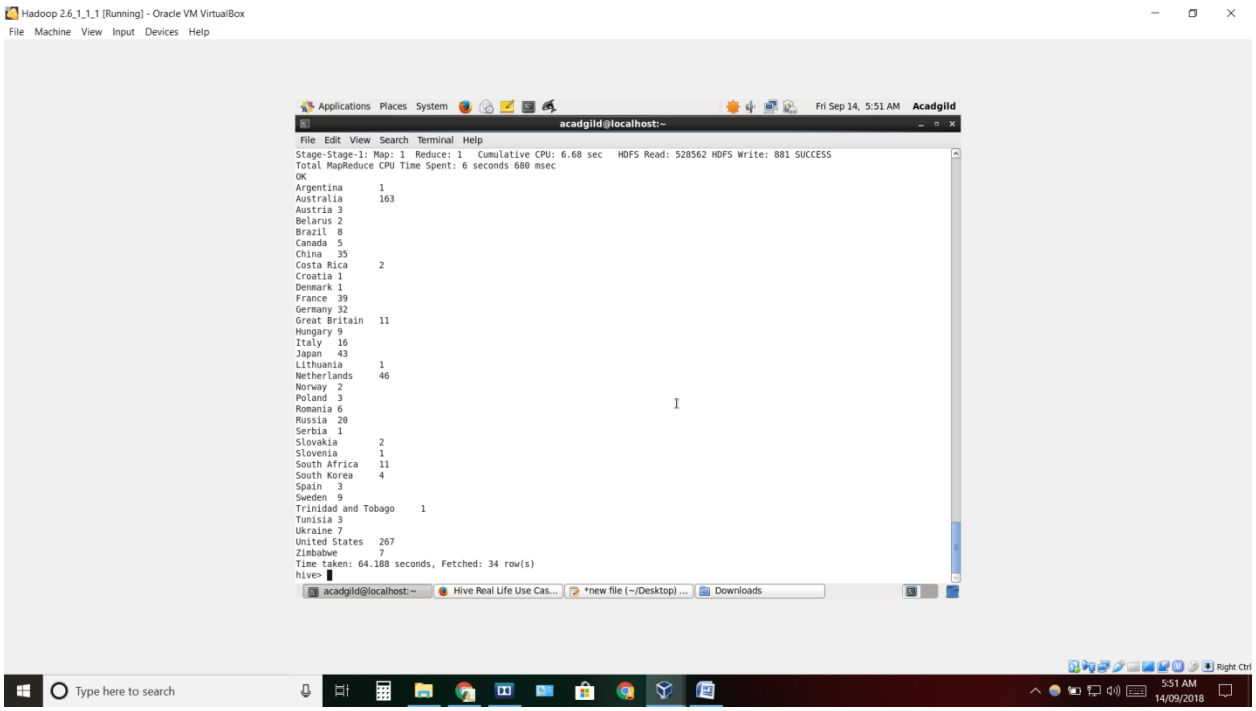

2. Write a Hive program to find the number of medals that India won year wise.
   select year,SUM(total) from olympic where country = "India" GROUP BY year;

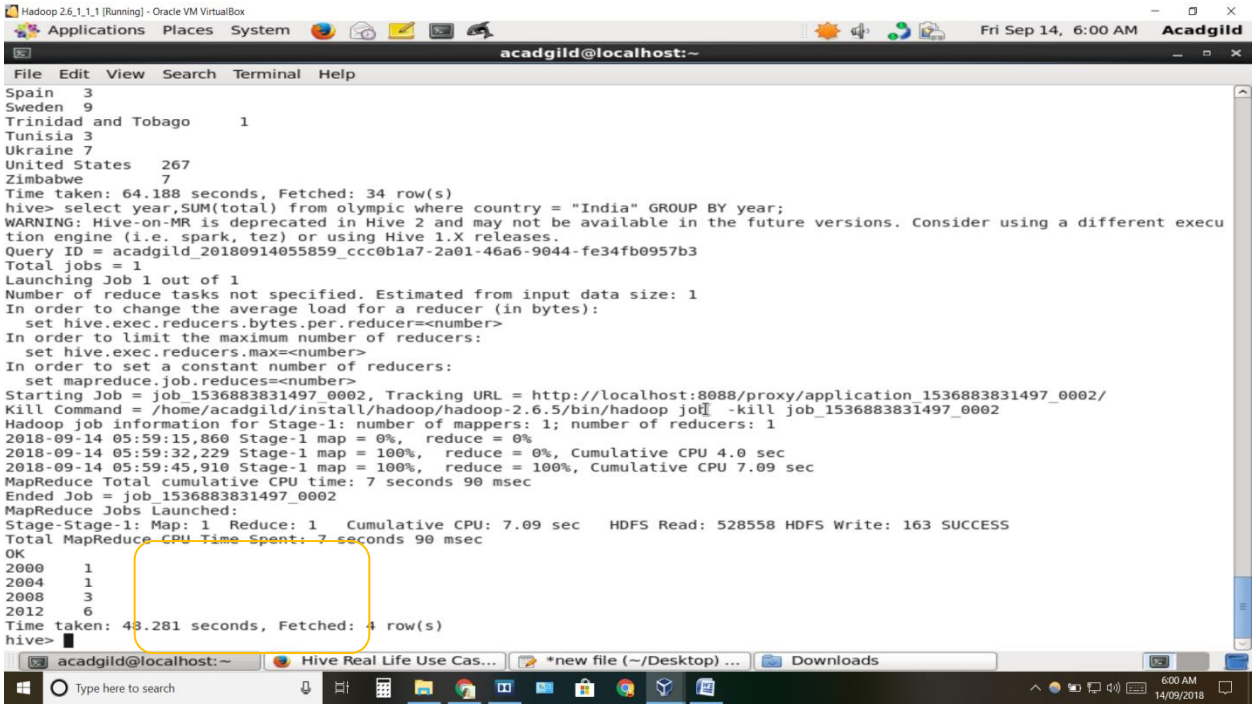

3. Write a Hive Program to find the total number of medals each country won.
   select country,SUM(total) from olympic GROUP BY country;

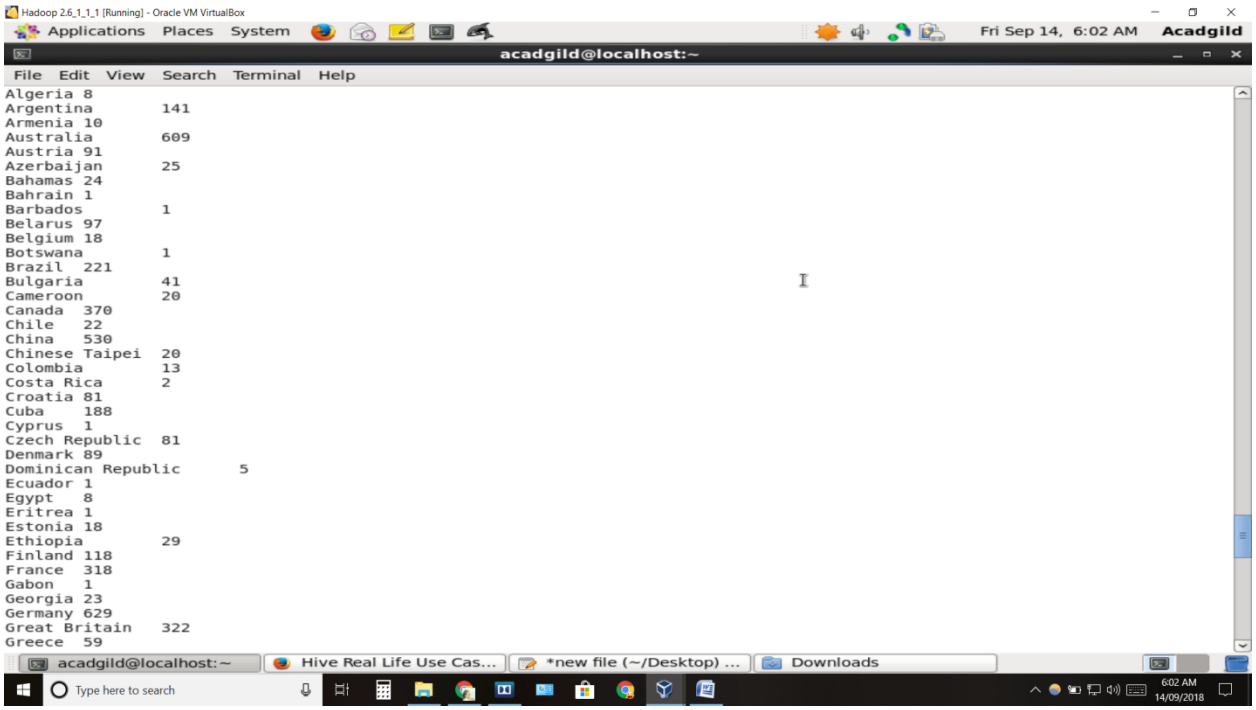

4. Write a Hive program to find the number of gold medals each country won.
   select country,SUM(gold) from olympic GROUP BY country;

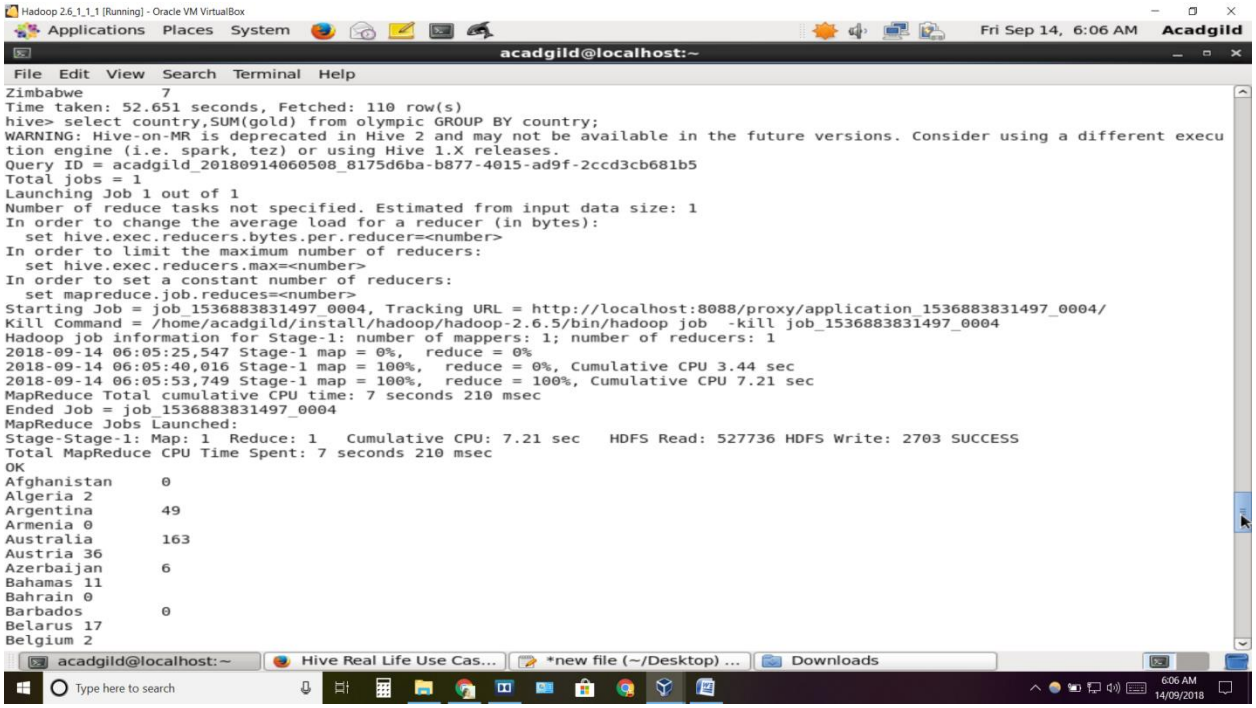

Task 2

Write a hive UDF that implements functionality of string concat_ws(string SEP, array<string>).

This UDF will accept two arguments, one string and one array of string.

It will return a single string where all the elements of the array are separated by the SEP.

Hive UDF

```
import java.util.ArrayList;
import org.apache.hadoop.hive.ql.exec.UDF;
public class JoinArray extends UDF{
public String evaluate (String separator, ArrayList<String> array) {
StringBuffer sBuffer;
if (array == null) {
return null;
}
sBuffer = new StringBuffer();
sBuffer.append(array.get(0));
for (int i=1; i < array.size(); i++) {
sBuffer.append(separator);
sBuffer.append(array.get(i));
}
return sBuffer.toString();
}
}
```

In below screen shot shows following

- Create table emparraynew;
- Describe table emparraynew;
- Load data from local in table emparrynew;
- add jar /home/acadgild/Desktop/empAr.jar;
- create temporary function separa as 'JoinArray.JoinArray';
- select separa(empName,empDesignation) from empArray;

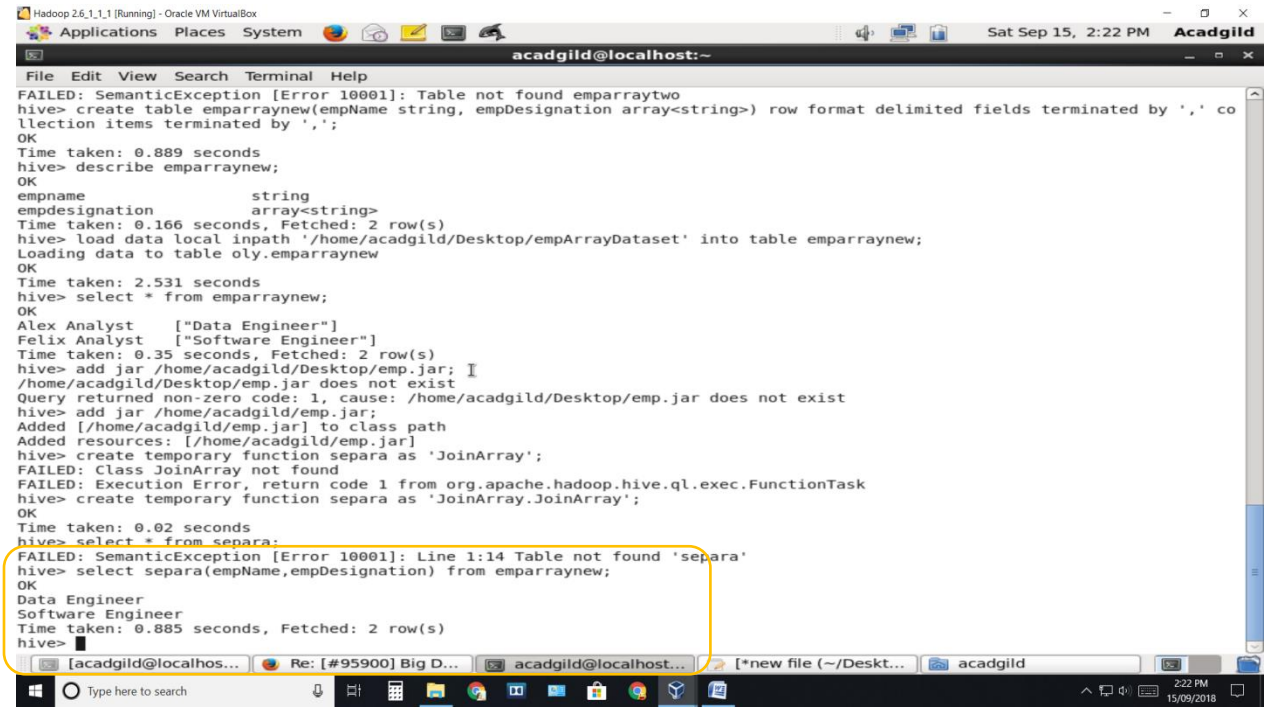

Task 3

Link: https://acadgild.com/blog/transactions-in-hive/

Refer the above given link for transactions in Hive and implement the operations given in the blog using your own sample data set and send us the screenshot.

Enabling ACID transactions in HIVE by using flowing tables.

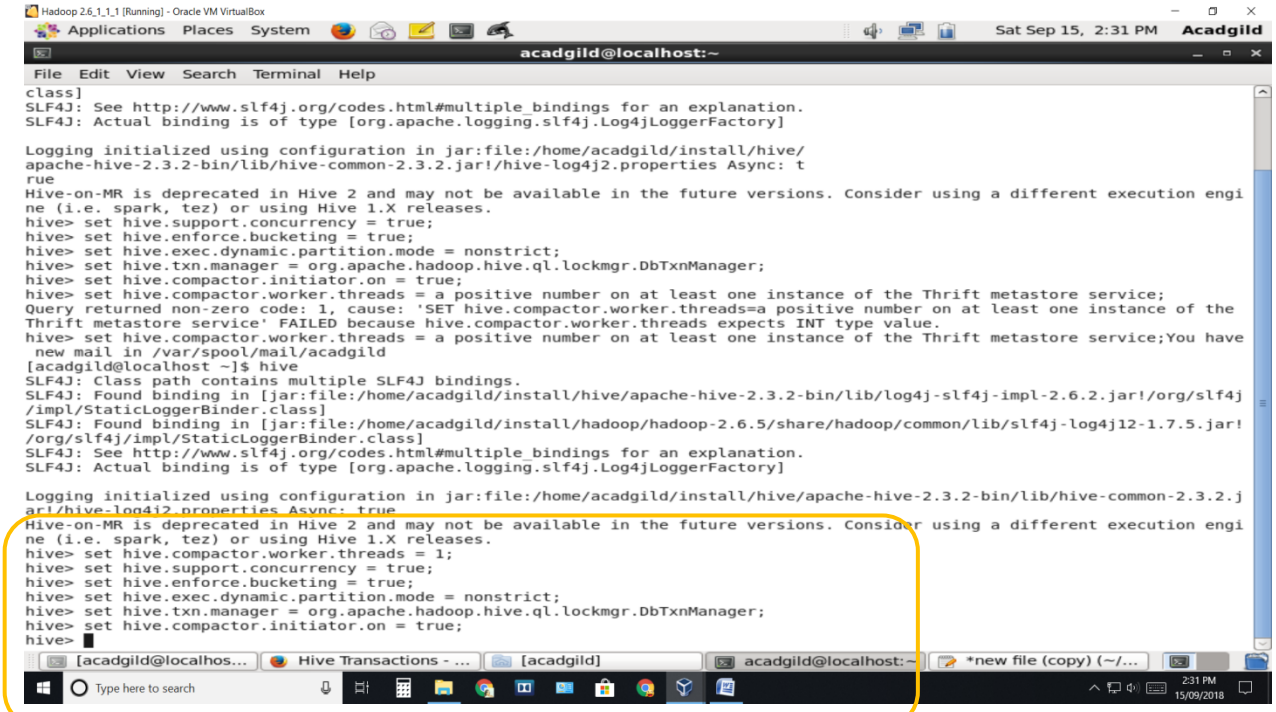

Create table and Insert data in table college;

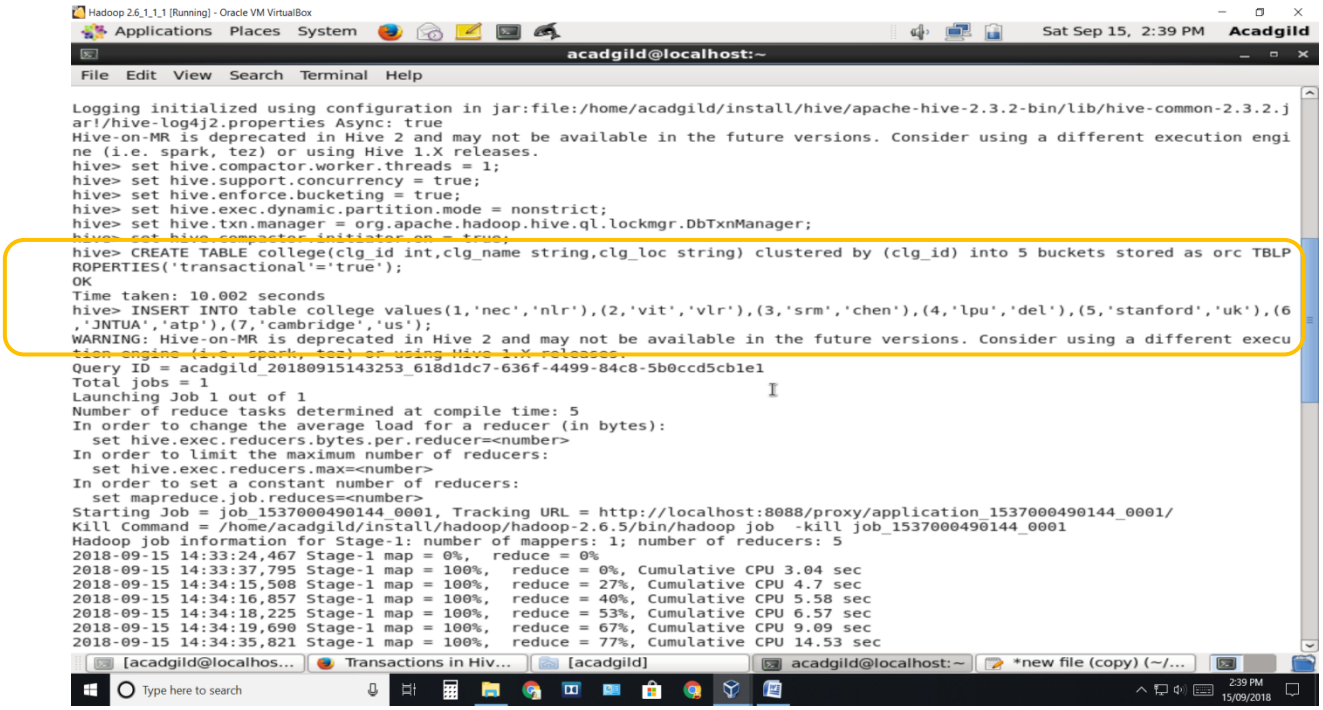

Check data in table after insertion

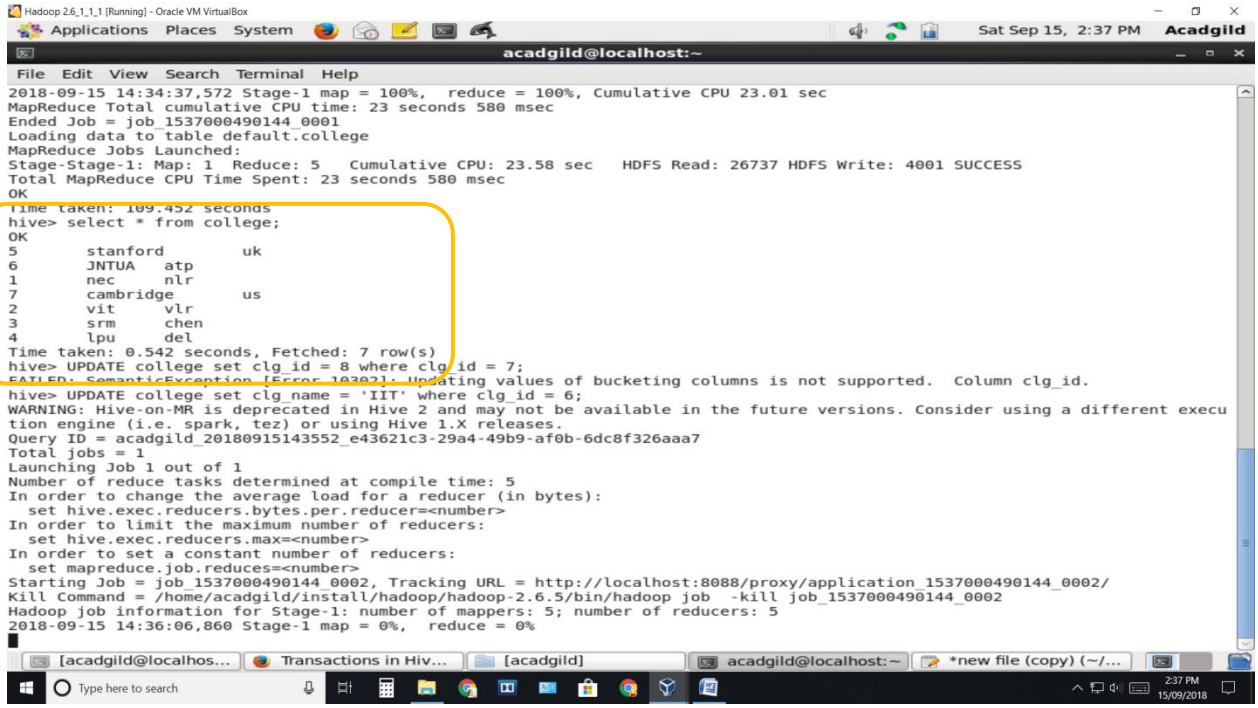

Updating the Data in Hive Table College

```
hive> UPDATE college set clg_name = 'IIT' where clg_id = 6;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execu
tion engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180915143552_e43621c3-29a4-49b9-af0b-6dc8f326aaa7
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537000490144_0002, Tracking URL = http://localhost:8088/proxy/application_1537000490144_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job  -kill job_1537000490144_0002
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-09-15 14:36:06,860 Stage-1 map = 0%,  reduce = 0%
2018-09-15 14:37:05,577 Stage-1 map = 60%,  reduce = 0%, Cumulative CPU 13.16 sec
2018-09-15 14:37:09,007 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 19.28 sec
2018-09-15 14:37:47,637 Stage-1 map = 100%,  reduce = 27%, Cumulative CPU 21.29 sec
2018-09-15 14:37:51,961 Stage-1 map = 100%,  reduce = 53%, Cumulative CPU 24.11 sec
2018-09-15 14:37:53,354 Stage-1 map = 100%,  reduce = 67%, Cumulative CPU 25.14 sec
2018-09-15 14:37:58,998 Stage-1 map = 100%,  reduce = 73%, Cumulative CPU 27.85 sec
2018-09-15 14:38:00,434 Stage-1 map = 100%,  reduce = 80%, Cumulative CPU 29.48 sec
2018-09-15 14:38:01,563 Stage-1 map = 100%,  reduce = 93%, Cumulative CPU 31.82 sec
2018-09-15 14:38:02,622 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 32.9 sec
MapReduce Total cumulative CPU time: 32 seconds 900 msec
```

```
Time taken: 133.148 seconds
hive> select * from college;
OK
5       stanford        uk
6       IIT     atp
1       nec     nlr
7       cambridge       us
2       vit     vlr
3       srm     chen
4       lpu     del
Time taken: 0.543 seconds, Fetched: 7 row(s)
hive>
```

[acadgild@localhos... Transactions in Hiv...

Deleting a Row from Hive Table College

```
hive> delete from college where clg_id=5;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the f
tion engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180915144718_dbc82677-e082-496e-a058-f6f6c6a6855e
```